

Communicating with Environment

Mainak Chaudhuri
Indian Institute of Technology Kanpur

Sketch

- Input and output (I/O)
- Interrupts
- Direct memory access (DMA)
- Exceptions
- System calls

2

I/O devices

- A computer that can only execute instructions for computing and accessing memory is not very useful
 - There is no way for the environment to give inputs to the computer or examine outputs of computation
- A typical computer interfaces with a large number of I/O devices
 - Keyboard, display, mouse, speaker, microphone, hard disk, printer, USB devices, etc.
 - Need a mechanism to communicate with these devices
 - How to read a key punch or how to display

3

I/O devices

- A typical I/O device has a set of command registers and a set of data registers
 - These are assigned unique addresses and can be read or written to through load or store instructions
 - Known as memory-mapped I/O registers
 - To the computer, these appear as memory locations
 - The only difference is that they are not in DRAM, but in I/O devices
 - For example, printing something on the printer involves storing the data to be printed to the printer data registers and an appropriate command to the printer command registers

4

I/O devices

- Certain I/O commands send responses back to the computer
 - A keyboard read needs to be conveyed to the computer
 - Completion of a disk read needs to be conveyed to the computer
 - Any read operation must be communicated to the computer
 - To detect completion of a disk or keyboard read, one possibility is to continuously poll a register of the disk controller or the keyboard controller
 - Wastes computer's time (computer could do something else during this time)

5

I/O devices

- Polling works only if the computer is aware that it will receive some response from a certain I/O device
 - Certain responses are accidental (e.g., ctrl+C to terminate a program or a mouse click)
 - In such cases, the computer did not know beforehand and was not polling the memory-mapped register
- An efficient solution that covers all cases is implemented using interrupts
 - These are signals sent by the I/O devices to the computer
 - For example, a key punch generates an interrupt

Interrupts

- Interrupts stop the normal instruction processing of a computer and make it execute an interrupt handler function
 - Interrupts can be generated by hardware (e.g., I/O devices) or software (e.g., exceptions and system calls)
 - MIPS treats all these as exceptions (any exceptional situation that interrupts normal instruction execution)
 - There are two ways to implement interrupt or exception handlers
 - Vectored interrupts or exceptions
 - Cause-based interrupt handling (non-vectored)

7

Vectored interrupts

- Reserve an area of memory (outside user memory map) to store interrupt handlers
 - Each interrupt is given a number and an array stores the starting addresses of the interrupt handlers
 - Starting address of interrupt handler i is stored in $IV[i]$
 - Assume that the interrupt number is in $\$1$ and the starting address of array IV is in $\$3$
- ```

sll $2, $1, 2
add $4, $3, $2
lw $4, 0($4)
jalr $4

```

8

### Vectored interrupts

- The handlers stored at vector locations are usually very small
  - These handlers further inspect the reason for the interrupt and jump to bigger interrupt service routines (ISRs)
    - For example, all hardware interrupts are usually given a single number and they all call the same interrupt handler
    - The interrupt handler finds out the source device of the hardware interrupt by consulting a special status register that stores the source/cause of interrupt
    - The interrupt handler calls the appropriate ISR

9

### Non-vectored interrupts

- MIPS implements non-vectored interrupts
  - All interrupts jump to the same fixed location (0x80000180)
  - MIPS maintains a status register and a cause register
    - Status register encodes the source of interrupts (six hardware and two software interrupt levels)
    - Cause register encodes the reason for interrupt (used mostly for software interrupts)
  - The interrupt handler examines the status and cause registers and jumps to the appropriate ISR
- ISR copies the input from I/O device data registers (e.g., keyboard buffer) to memory

### Direct memory access (DMA)

- Copying large amounts of data from input devices to memory may take a significant amount of time
  - Reading files from disk
- Same applies to copying from memory to output devices
  - Write to files on disk
- Occupying the computer during this time wastes computer's resources
  - Could compute something useful during this time
- Direct memory access (DMA) frees up the computer during data copying from/to I/O

11

### Direct memory access (DMA)

- Computer initializes a specialized hardware called DMA controller by setting up the copy address range and the number of bytes to be copied
- DMA controller does the actual copy operation
  - DMA controller when copying from an input device sends the data to the DRAM controller for writing
  - DMA controller when copying to an output device sends read requests to the DRAM controller
  - When the copying completes, the DMA controller sends an interrupt to the computer to notify about the DMA completion

12

## Exceptions

- Exceptions refer to situations where the running program exhibits unexpected behavior
  - Arithmetic overflow, divide by zero, fetching and decoding an illegal opcode, accessing an illegitimate address or unaligned address (MIPS)
  - In such situations, the PC of the offending instruction is saved in a register called exception PC (EPC) and the program counter is changed to point to a location that has a “trap” instruction
  - The trap instruction allows the computer to enter the operating system which further invokes the appropriate exception handler after examining the cause register

13

## Exceptions

- Some exceptions are restartable while some exceptions are not
  - Restartable exceptions return to normal execution of the program by copying the EPC or EPC+4 (depending on the situation) into a general-purpose register (\$X) and issuing a jr \$X instruction after the exception handler completes
    - Example: non-availability of code/data in memory (causes a page fault exception)
  - Non-restartable exceptions typically lead to termination of the running program with an appropriate message printed on display
    - Arithmetic exceptions, illegal opcode, crossing legitimate memory boundary, etc.

14

## System calls

- System calls are pseudo-function calls to request access to certain hardware/software resources of the computer system
  - Reading from a file on disk, reading from keyboard, writing to display, writing to a file, allocating dynamic memory, etc. involve system calls
  - Some computers refer to system calls as software interrupts
  - MIPS ISA has the syscall instruction for this purpose
    - R format, has no operand, function 0xc
    - Before invoking the syscall instruction, few registers need to be set up with appropriate information

15

## syscall instruction

- Every system call has a number to indicate the purpose of the system call
  - Reading from a file, writing to a file, allocating dynamic memory all have different system call numbers
  - The system call number must be placed in register \$v0 in MIPS
  - A system call can accept four arguments in registers \$a0, \$a1, \$a2, \$a3
    - More than four arguments can be passed by packing fourth argument onward in a structure and passing a pointer to the structure in \$a3
  - A system call returns any expected response in \$v0

16

### syscall instruction

- A syscall instruction executes by first changing the program counter to jump to a location
  - Vectored implementations treat all system calls as software interrupts and all system calls are assigned a fixed location in interrupt vector array
  - Non-vectored implementations jump to the same location for all kinds of interrupts
  - The code in this location further examines \$v0 and jumps to the appropriate system call handler
  - System call handler accesses the argument registers and does the necessary things

17

### Examples of system call

- Consider the read system call of UNIX
  - Invoked as part of C library functions scanf, fscanf, read, etc.
    - stdin is treated like a file
  - Three arguments: \$a0 should have file descriptor (a non-negative integer representing the file), \$a1 should have destination memory buffer address, \$a2 should have number of bytes to read
    - Number of bytes to read is inferred from the data types of the arguments of the high-level function call
  - SPIM simulator uses different system calls and conventions for reading from keyboard
    - Only read\_string requires buffer address and length in \$a0 and \$a1; other keyboard reads do not have args

### Examples of system call

- Consider the read system call
  - On return of a read system call, \$v0 contains the number of bytes read
    - Note that the actual bytes read can be found in memory starting from the address passed in \$a1 to the system call
  - SPIM uses a different convention for reading from keyboard: on return of a read system call for reading from keyboard, \$v0 or \$f0 (depending on type of data read) contains the actual value read (except for read\_string, which does not have any return value)

19

### Examples of system call

- Complete path of scanf
  - C program calls scanf function
  - Program jumps to scanf library function (jal)
  - Sets up \$v0, \$a0, \$a1, \$a2
  - syscall instruction executes (part of scanf func.)
  - Invokes system call handler for reading from file
  - Syscall handler goes to sleep until interrupted
  - Keyboard punch generates interrupt to computer
  - ISR wakes up system call handler
  - Tail code of system call handler copies data from keyboard buffer to computer memory pointed to by \$a1
  - Returns to C program that called scanf

20

### Examples of system call

- Consider the write system call
  - Invoked as part of C library functions printf, fprintf, write, etc.
    - stdout is treated like a file
  - Three arguments: \$a0 should have file descriptor (a non-negative integer representing the file), \$a1 should have source memory buffer address containing the characters to be written, \$a2 should have number of bytes to write
    - Number of bytes to write is inferred from the data types of the arguments of the high-level function call
  - SPIM simulator uses different system calls and conventions for writing to display
    - Only argument is \$a0 or \$f12 holding the value to be printed (\$a0 contains a pointer to string in print\_string)

### Examples of system call

- Complete path of printf
  - C program calls printf function
  - Program jumps to printf library function (jal)
  - Sets up \$v0, \$a0, \$a1, \$a2
  - syscall instruction executes (part of printf func.)
  - Invokes system call handler for writing to file
  - Copies data (\$a2 bytes) from source buffer pointed to by \$a1 to the display device's memory
  - Invokes the display device driver to actuate the low-level devices (CRT or LCD or LED) for printing the characters
  - Returns to C program that called printf

22

### Examples of system call

- Consider the sbrk system call used for allocating dynamic memory
  - Invoked as part of C library function malloc
  - One argument: \$a0 should contain the number of bytes to be allocated
  - On return, \$v0 contains the starting address of the allocated memory region

23

### Example of syscall in SPIM

- SPIM assembly language program for printing the string "Hello World\n" to display

```
.data
msg: .asciiz "Hello World\n"

.text
.globl main
main: li $v0, 4 # syscall 4 (print_str)
 la $a0, msg # argument: string address
 syscall # print the string

 jr $ra # return to caller
```

24

### Example of syscall in SPIM

- SPIM assembly language program for reading an integer, adding 42 to it, and printing the result to display

```
.text
.globl main
main: li $v0, 5 # syscall 5 (read_int)
 syscall
 addi $a0, $v0, 42 # print_int argument
 li $v0, 1 # syscall 1 (print_int)
 syscall # print the integer

 jr $ra # retrun to caller
```

25

### Example of syscall in SPIM

- SPIM assembly language program for reading two floats, adding them, and printing the result to display

```
.text
.globl main
main: li $v0, 6 # syscall 6 (read_float)
 syscall
 mov.s $f12, $f0 # return value in $f0
 li $v0, 6
 syscall
 add.s $f12, $f0, $f12 # print_float arg
 li $v0, 2 # syscall 2 (print_float)
 syscall # print
 jr $ra # retrun to caller
```

26

### Example of syscall in SPIM

- SPIM assembly language program for reading ten integers, storing them in an array, adding them, and printing the result

```
.data
arrayX: .space 40
msg: .asciiz "Sum of values: "
endmsg: .asciiz "\n"

.text
.globl main
main: addi $t0, $0, 10 # i ← 10
 la $t1, arrayX # $t1 ← arrayX
loop: li $v0, 5 # syscall 5 (read_int)
 syscall
 sw $v0, 0($t1) # *arrayX ← $v0
 addi $t1, $t1, 4 # arrayX++
 addi $t0, $t0, -1 # i--
 bne $t0, $0, loop

 jr $ra # retrun to caller
```

27

### Example of syscall in SPIM

```
 addi $t0, $0, 10 # i ← 10
 la $t1, arrayX # $t1 ← arrayX
 xor $t2, $t2, $t2 # sum ← 0
loop1: lw $v0, 0($t1) # $v0 ← *arrayX
 add $t2, $t2, $v0 # sum ← sum + $v0
 addi $t1, $t1, 4 # arrayX++
 addi $t0, $t0, -1 # i--
 bne $t0, $0, loop1

 li $v0, 4 # print_string
 la $a0, msg # argument: string
 syscall # print the string
 li $v0, 1 # print_int
 add $a0, $t2, $0 # int to print
 syscall

 li $v0, 4 # print_string
 la $a0, endmsg # argument: string
 syscall # print the string

 jr $ra # retrun to caller
```

28

### System calls supported by SPIM

- print\_int: no. 1, \$a0 should have int value
- print\_float: no. 2, \$f12 should have float value
- print\_double: no. 3, (\$f12, \$f13) should have double value
- print\_string: no. 4, \$a0 should have the pointer to the string
- print\_char: no. 11, \$a0 should have the ASCII value of char
- sbrk: no. 9, \$a0 should have the number of bytes to be allocated, return address in \$v0

### System calls supported by SPIM

- read\_int: no. 5, return int value in \$v0
- read\_float: no. 6, return float value in \$f0
- read\_double: no. 7, return double value in (\$f0, \$f1)
- read\_string: no. 8, \$a0 should have the destination memory buffer address, \$a1 should have the length of the string
- read\_char: no. 12, return char value in \$v0
- exit: no. 10 (terminates the calling program)
- exit2: no. 17 (terminates spim)

30

### System calls supported by SPIM

- Also support system calls to open a file, read from a file, and write to a file
  - Useful for programs operating on files

31