

Processor FSM Design

Mainak Chaudhuri
Indian Institute of Technology Kanpur

Sketch

- Abstract model of computer
- Single-cycle instruction execution
- Multi-cycle instruction execution
- Pipelined instruction execution

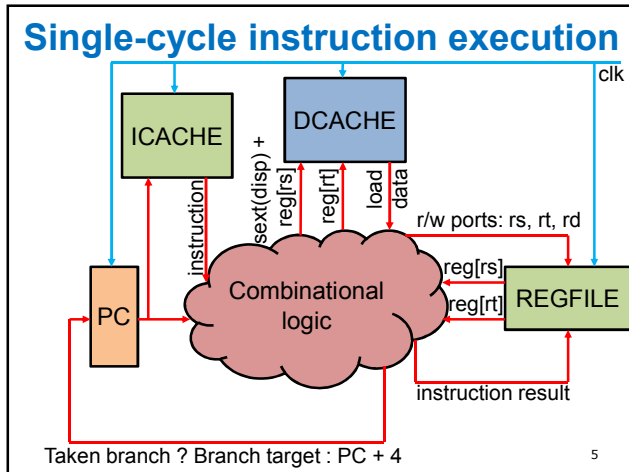
2

Abstract model of computer

- Each instruction undergoes five stages
 - Stage 0 (IF): fetch the instruction pointed to by program counter from memory
 - Stage 1 (ID/RF): decode the instruction to extract various fields and read source register operands
 - Source registers can be read in parallel with instruction decoding in MIPS due to fixed positions of rs and rt in all formats of the ISA
 - Stage 2 (EX): execute the instruction in ALU; compute address of load/store instructions; update program counter to PC+4 or branch target (if this instruction is control transfer ins)
 - Stage 3 (MEM): access memory if load/store instruction; use address computed in stage 2
 - Stage 4 (WB): write result back to destination register if the instruction produces a result ³

Single-cycle instruction execution

- Possible to implement all five stages as a big combinational logic
 - The logic will still access a few sequential logic elements e.g., program counter register (stages 0 and 2), register file (in stages 1 and 4), icache (in stage 0), and dcache (in stage 3)
 - These can be viewed as inputs and outputs of the combinational logic
 - PC, icache[PC], reg[rs], reg[rt], dcache[disp(rs)] are inputs
 - » dcache[disp(rs)] is a relevant input for loads
 - PC, reg[rd] or reg[rt], dcache[disp(rs)] are outputs
 - » mem[disp(rs)] is a relevant output for stores
 - The instruction opcode and function (for R format) fields decide the control of combinational logic e.g., which register operand(s) is/are valid sources, which ALU operation to invoke, etc. ⁴



Single-cycle instruction execution

- Taken branch

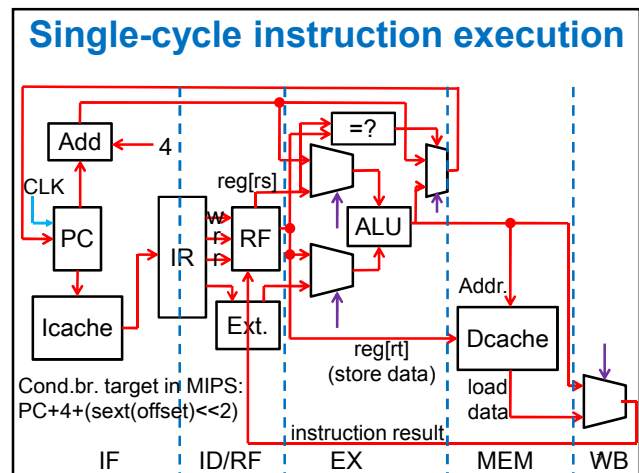
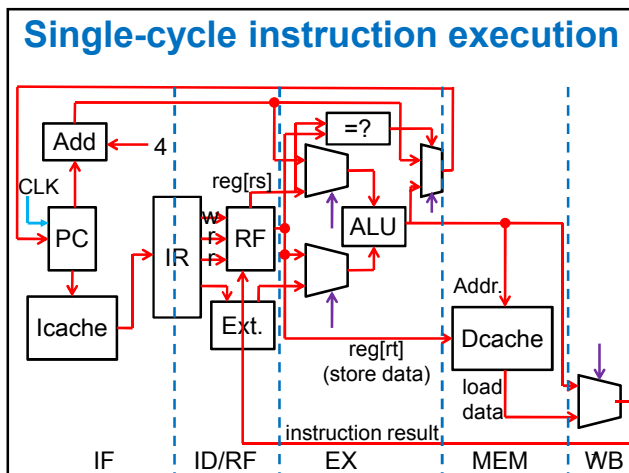
- A branch instruction whose condition dictates that the next instruction to be executed is the one at the branch target

- beq \$1, \$2, label (the branch is taken if $\$1 == \2)
- Unconditional direct and indirect jumps and procedure calls are all taken branches

- Not taken branch

- A branch instruction whose condition dictates that the next instruction to be executed is the instruction right next to the branch

- A la Frost's "Road Not Taken"
- beq \$1, \$2, label (the branch is not taken if $\$1 != \2)
- A conditional branch can be taken or not taken depending on how the condition evaluates at run time



Single-cycle instruction execution

- Clock cycle time must be large enough to accommodate the lengthiest instruction
 - This is typically the load instruction whose critical path includes all five stages: IF: fetch instruction, ID/RF: decode while reading reg[rs] and sign extending displacement, EX: add reg[rs] with sext(displacement) to generate address, MEM: read dcache to get load data, WB: write load data back to reg[rt]
 - Not a good idea since many other instructions take smaller amount of time
 - Control transfer instructions require only three stages
 - Stores require only four stages
 - All ALU instructions can bypass MEM stage

8

Multi-cycle instruction execution

- Each stage takes one cycle to execute
 - Processor is now a five state FSM
 - Need flip-flops at stage boundaries and each stage is a combinational logic
 - Clock cycle time = latency of the longest stage
 - Most instructions take five cycles
 - ALU instructions do nothing in MEM stage, but will have to go through it if FSM state transition is implemented using an incrementer
 - Some instructions may take longer e.g., multiply/divide, load/store (in case of dcache miss)
 - Some instructions may take shorter time e.g., stores take four cycles, control transfer instructions take three cycles
 - Overall CPI depends on instruction mix

9

Multi-cycle instruction execution

- Example
 - IF: 2 ns, ID/RF: 1 ns, EX: 1 ns, DMEM: 3 ns, WB: 1 ns
 - Branch frequency: 20%
 - Store frequency: 10%
 - Multiply/divide frequency: 5%, latency: 30 ns
 - Total instruction count: 100
- Multi-cycle
 - Cycle time: 3 ns, frequency: 333 MHz
 - CPI: $0.2*4 + 0.1*4 + 0.05*10 + 0.65*5 = 4.95$
 - Execution time: $100*4.95*3 \text{ ns} = 1485 \text{ ns}$
- Single-cycle
 - Cycle time: 8 ns, frequency: 125 MHz
 - CPI: $1*0.95 + \text{ceil}(30/8)*0.05 = 1.15$
 - Execution time: $100*1.15*8 \text{ ns} = 920 \text{ ns}$

10

Multi-cycle instruction execution

- Same example with a balanced pipeline
 - IF: 2 ns, ID/RF: 2 ns, EX: 2 ns, DMEM: 2 ns, WB: 2 ns
 - Branch frequency: 20%
 - Store frequency: 10%
 - Multiply/divide frequency: 5%, latency: 30 ns
 - Total instruction count: 100
- Multi-cycle
 - Cycle time: 2 ns, frequency: 500 MHz
 - CPI: $0.2*4 + 0.1*4 + 0.05*15 + 0.65*5 = 5.2$
 - Execution time: $100*5.2*2 \text{ ns} = 1040 \text{ ns}$
- Single-cycle
 - Cycle time: 10 ns, frequency: 100 MHz
 - CPI: $1*0.95 + \text{ceil}(30/10)*0.05 = 1.1$
 - Execution time: $100*1.1*10 \text{ ns} = 1100 \text{ ns}$ (worse than multi-cycle)

11


Pipelined instruction execution

- Observations from multi-cycle design
 - In the second cycle, I know if it is a branch; if not, start fetching the next instruction?
 - When the ALU is doing an addition (say), the decoder is sitting idle; can we use it for some other instruction?
 - In summary, exactly one stage is active at any point in time: wastes hardware resources
- Form a pipeline
 - Process five instructions in parallel
 - Each instruction is in a different stage of processing (called pipe stage)

12

Pipelined instruction execution

- Individual instruction latency is five cycles, but ideally can finish one instruction every cycle after the pipeline is filled up
 - Ideal CPI of 1.0 at the clock frequency of multi-cycle design
 - Execution time is ideally one-fifth of the multi-cycle design
 - Instruction throughput improves five times (number of instructions completed in a given time)



Pipelined instruction execution

- Gains
 - Extracting parallelism from a sequential instruction stream: known as instruction-level parallelism (ILP)
 - Can complete one instruction every cycle (ideally)
- Each pipe stage may get lengthened a little bit due to control overhead (skew time, setup time, propagation delay): limits the gain due to pipelining
 - Each instruction may take slightly longer for this reason

14