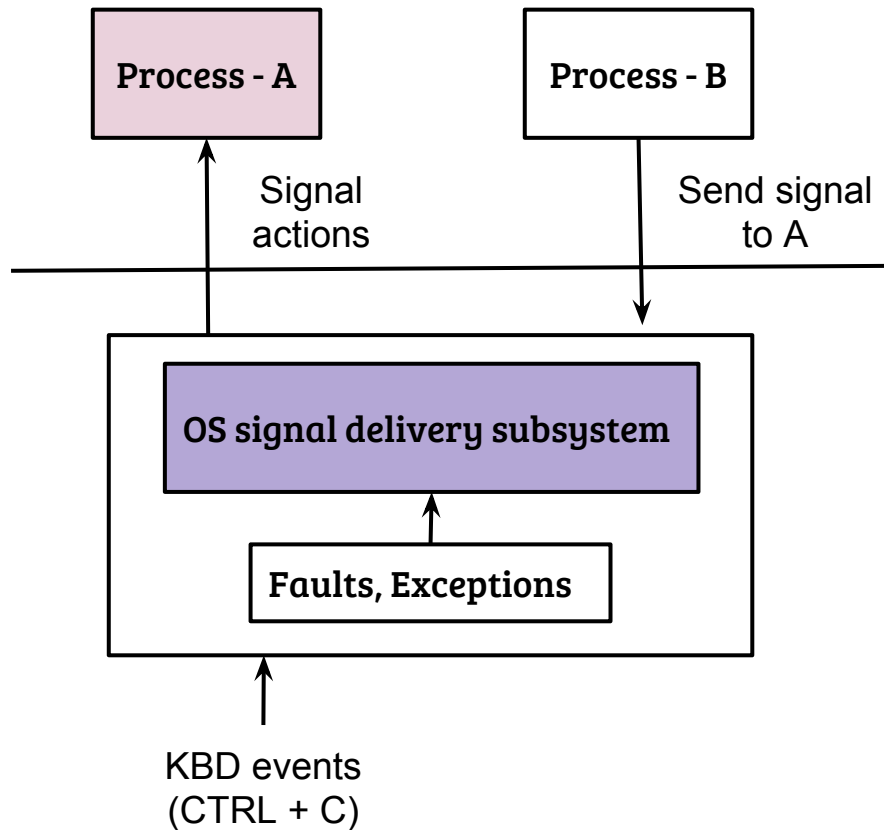


Operating Systems

Signals, Shell operations and IPC techniques

Debadatta Mishra, CSE, IITK

Why Signals?




- Termination of a process
- Process induced exceptions, div-by-zero, access to illegal memory area etc.
- Notification to process, SIGALRM, SIGCHLD
- Interprocess communication
- Custom actions on events

Signal semantics

Destination process

signal (signum, handler)



SIGHUP
SIGINT
SIGALRM
SIGUSR1
SIGCHLD

.....

1. SIG_IGN: Ignore the signal
2. SIG_DFL: Default action
3. Function address: custom handling

- If signal handler not registered, process is terminated (mostly)
- SIGKILL and SIGSTOP → no custom actions
- How does the handler invoked?

Signal semantics

Source process

kill (pid, signum)

**PID Of
target
process**

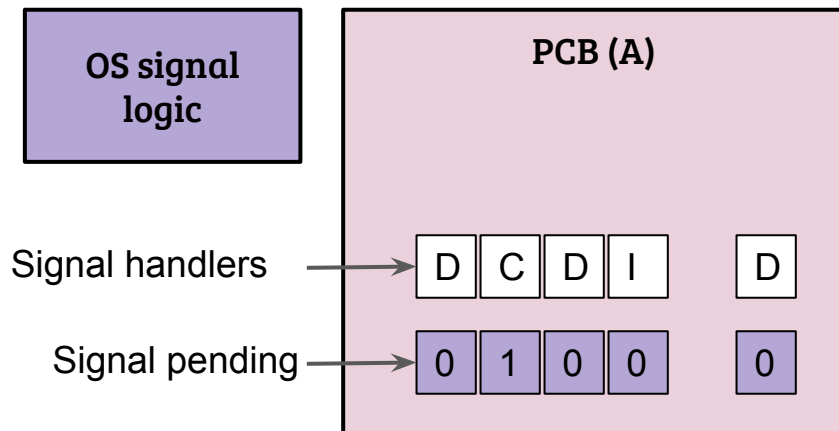
**SIGHUP
SIGINT
SIGALRM
SIGUSR1
SIGCHLD
.....**

- If pid == 0, signal is sent to all processes in the process group
- Must have permissions to send signals → same user or root user

OS support for signals

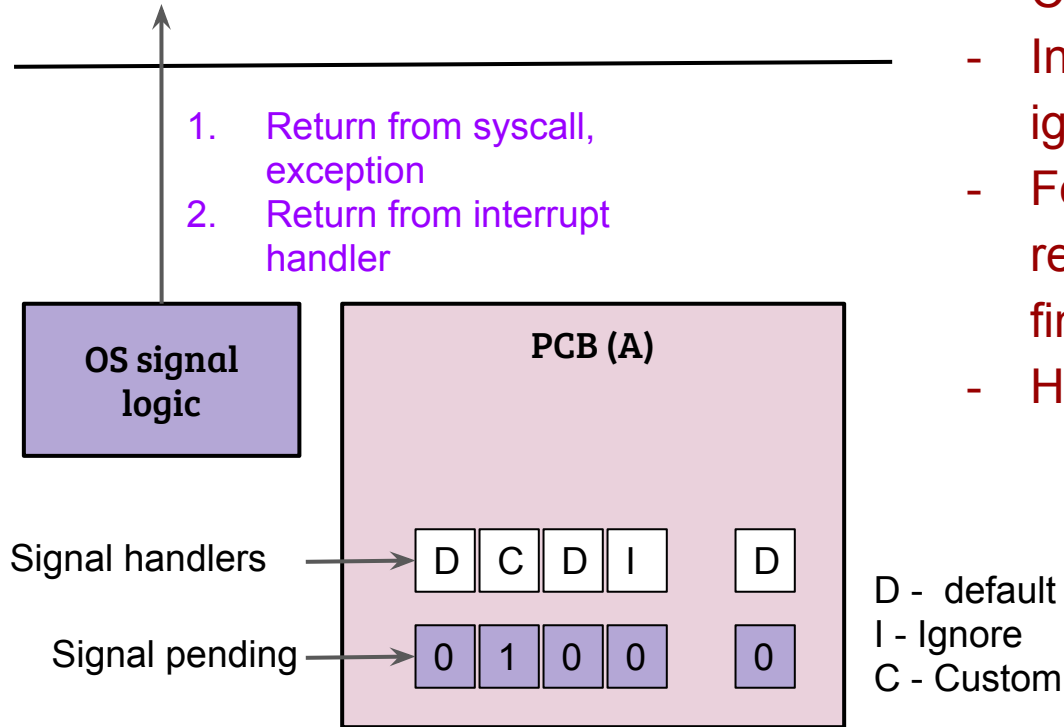
```
Process - A
main(){
    .....
    signal(1, H);
    .....
}
H(){
    .....
}
```

- Signal pending is modified by either kill () or OS
- Signal handlers are registered by the OS or the process
- Synchronous delivery is easy but not always possible
- Asynchronous delivery is complex



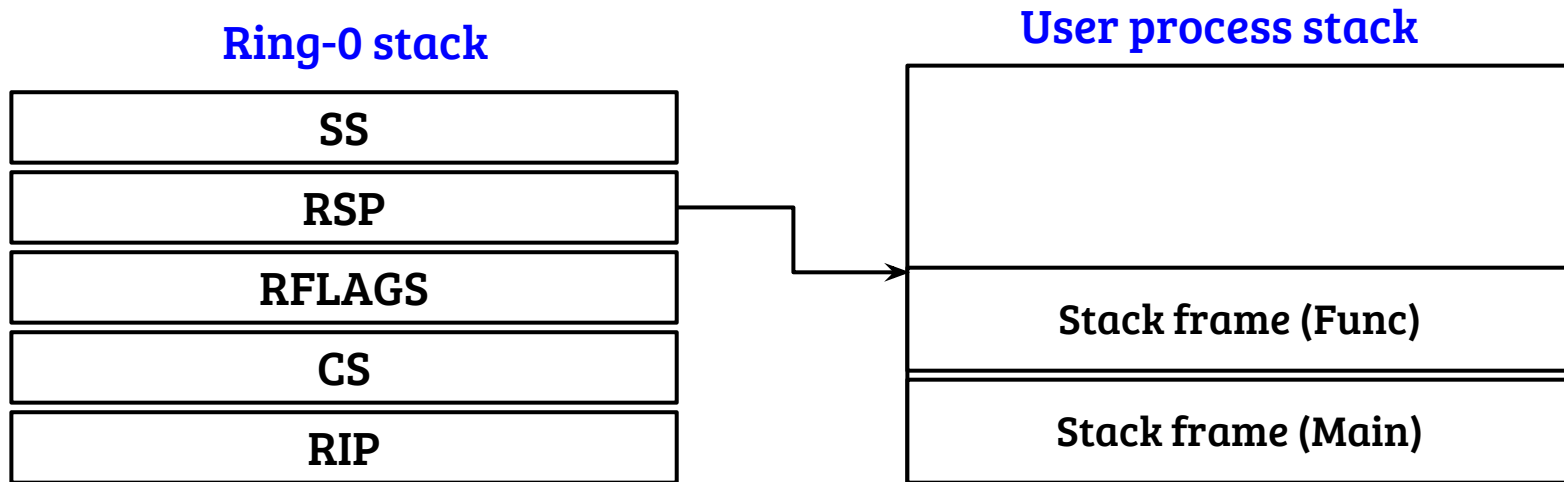
D - default
I - Ignore
C - Custom

When to deliver signals?



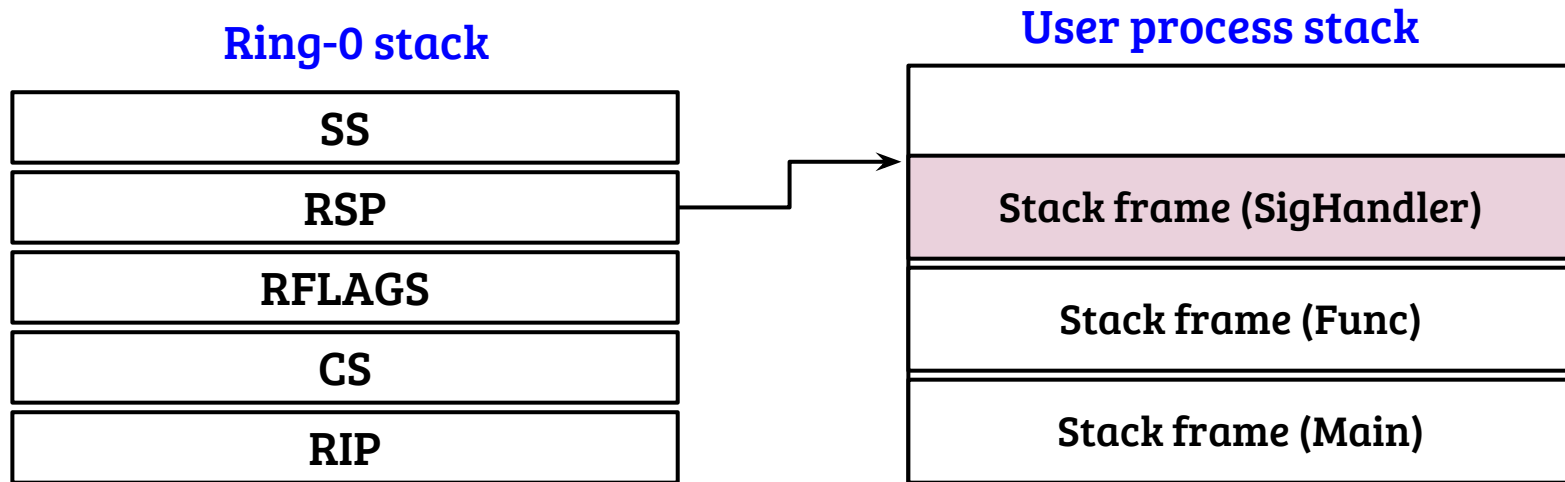
- Check for signal pending bit
- Invoke handler actions → default, ignore or custom
- For custom handlers, execution resumes after the handler is finished.
- How?

Signal delivery



- Assume that the program was interrupted when it was executing “Func”
- The OS can modify the user stack and the RIP. What values?

Signal delivery

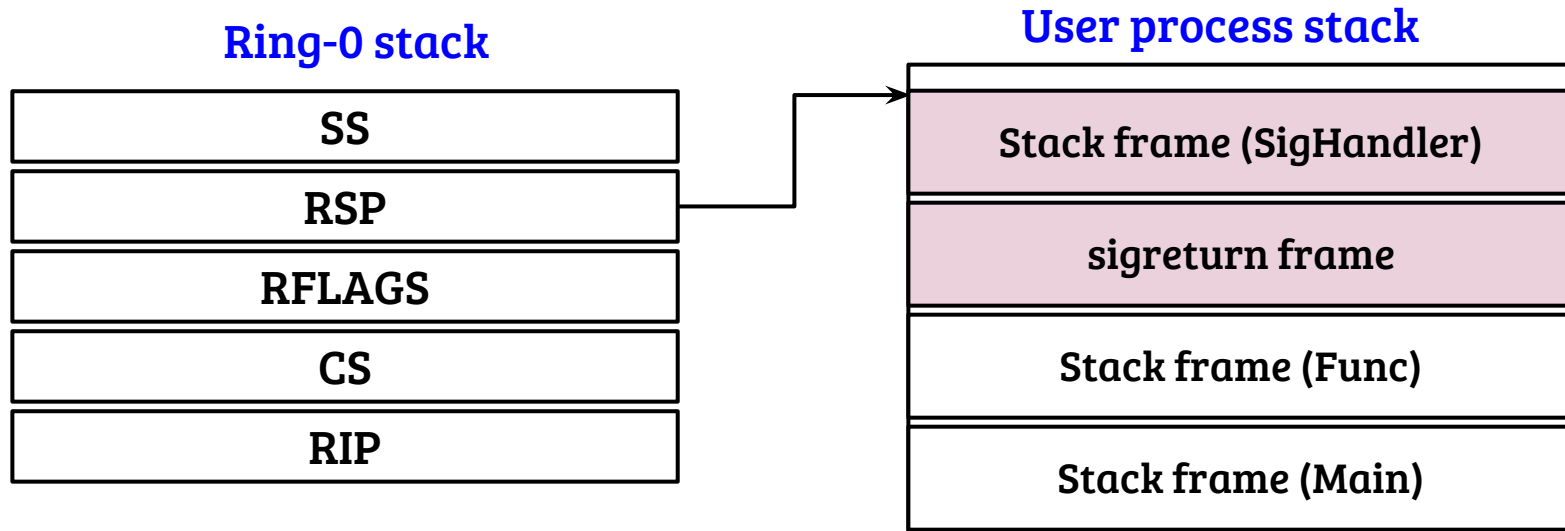


- Assume that the program was interrupted when it was executing “Func”
- The OS can modify the user stack and “mimic” a function call

Signal handling: design choices

- Remove handler → Handle signal → Register signal handler again
- Don't remove handler → New signal during handling → invoke the handler
- Temporarily disable signal → handle signal → re-enable signal

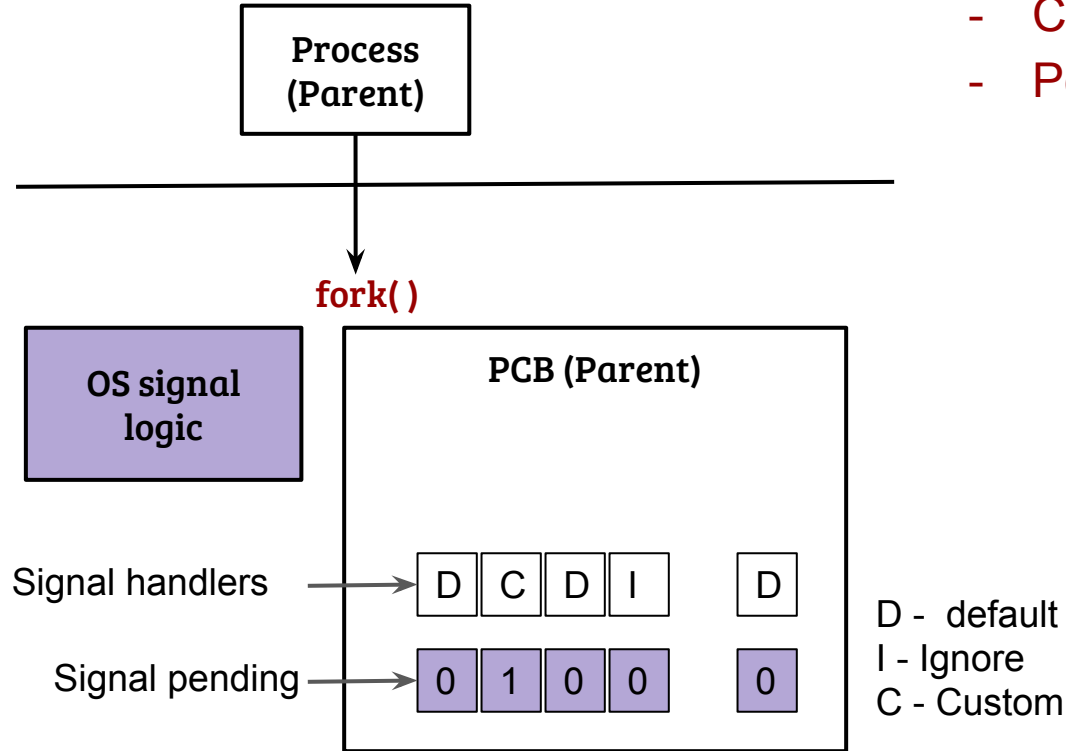
Signal delivery with handler notification (Linux)



- Sigreturn function invokes a syscall to reenale signal
- OS restores the stack to original state

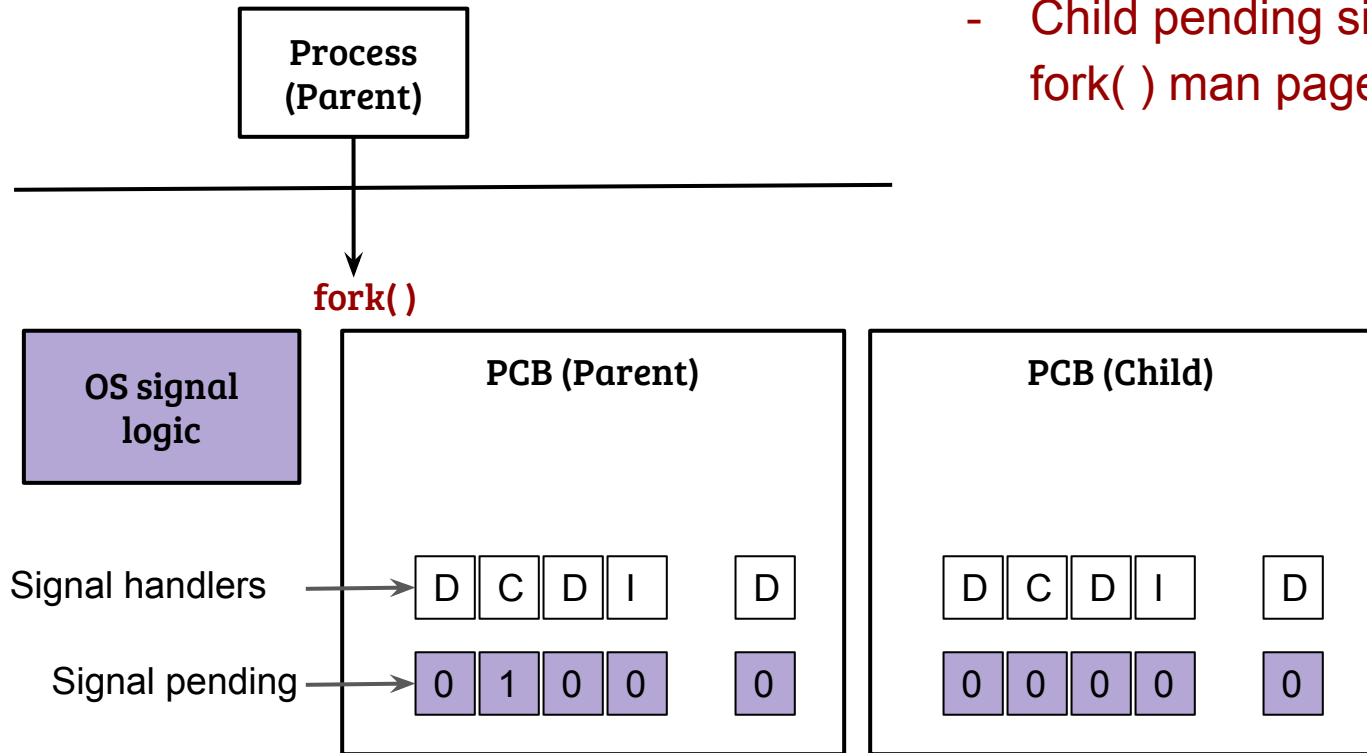
Signal and fork()

- Child inherits the signal handlers
- Pending signals?

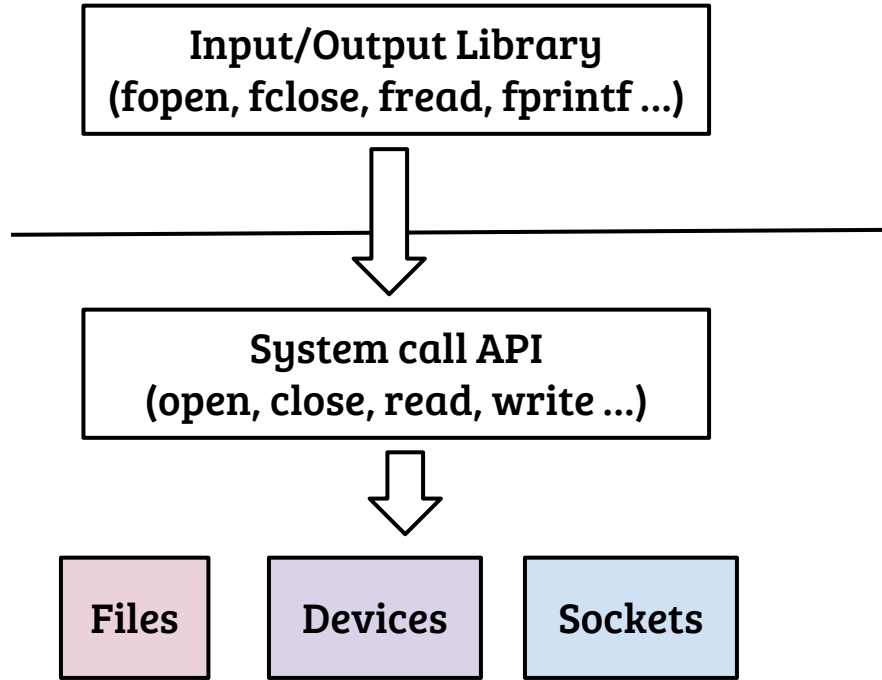


Signal and fork()

- Child pending signals is empty -- fork() man page. Why?



System calls (open, close, read, write)



- User process identify files through a file handle a.k.a. file descriptors
- In UNIX, the POSIX file API is used to access files, devices, sockets etc.


File system calls: getting a handle

Open: Provides a handle to a file

```
int open (char *path, int flags, mode_t mode)
```

- Access mode specified in flags : O_RDONLY, O_RDWR, O_WRONLY
- Access permissions check
- On success, a file descriptor (int) is returned
- If flags contain O_CREAT, mode specifies the file creation mode
- Open fds remain open across exec(), can be changed by setting O_CLOEXEC flag

Process view of file



- Per-process file descriptor table with pointer to "file" object
- `fd` \rightarrow file (many-to-one)
- file \rightarrow inode (many-to-one)
- On `fork()`, child inherits open file handles
- 0, 1, 2 are STDIN, STDOUT and STDERR, respectively

Read and Write

`ssize_t read (int fd, void *buf, size_t count);`

- `fd` → file handle
- `buf` → user buffer as read destination
- `count` → #of bytes to read



`read ()` returns #of bytes actually read, can be smaller than `count`

`ssize_t write (int fd, void *buf, size_t count);`

- Similar to `read`

Duplicate file handles (dup and dup2)


Before dup()



- Lowest available duplicate file handle is returned
- Example:

After dup()

+ dup(fd1)



- `close(1); dup(fd)`
- `dup2(old, new)` performs both steps in one system call

Shell redirection

- Example: `ls > tmp.txt`
- How implemented?

Shell redirection

- Example: `ls > tmp.txt`
 - How implemented?
-
- `fd = open ("tmp.txt")`
 - `close(1); close(2);` `// close STDOUT and STDERR`
 - `dup(fd); dup(fd)` `// 1 → fd, 2 → fd`
 - Invoke `exec()`

Shell: ls | wc -l

- Output of “ls” is input to “wc -l”
- How implemented?


- Option 1

temporary intermediate file + dup(0)

- Option 2

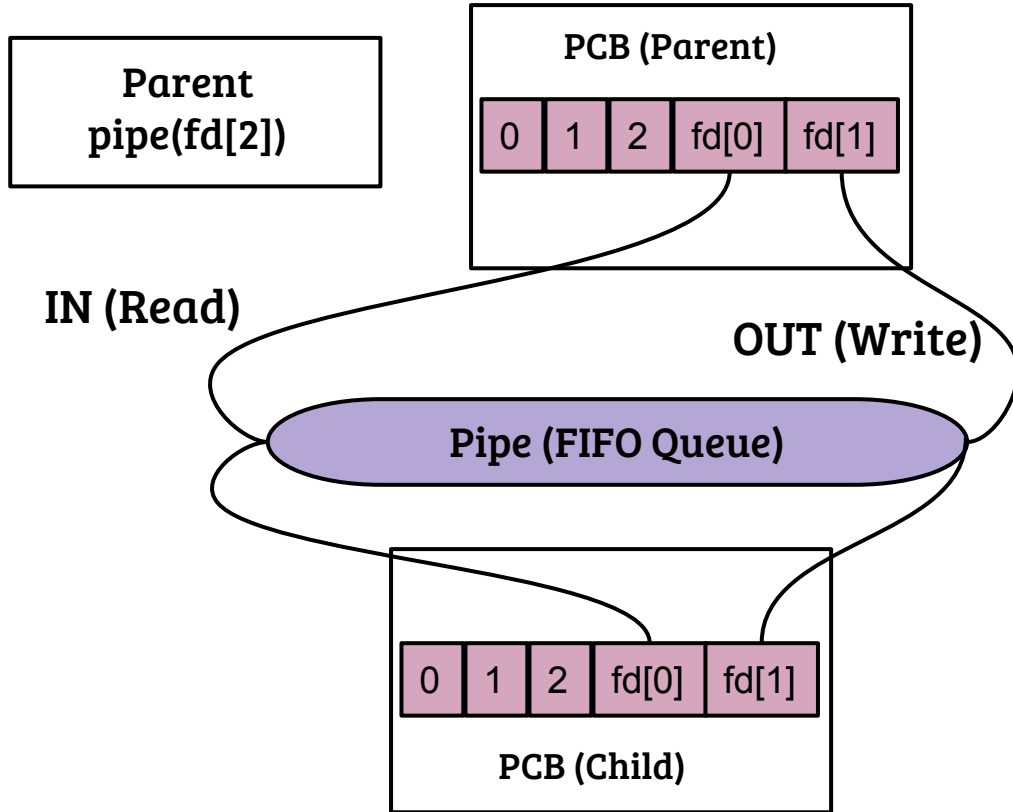
Shared buffer between “ls” and “wc” process

UNIX pipe() system call



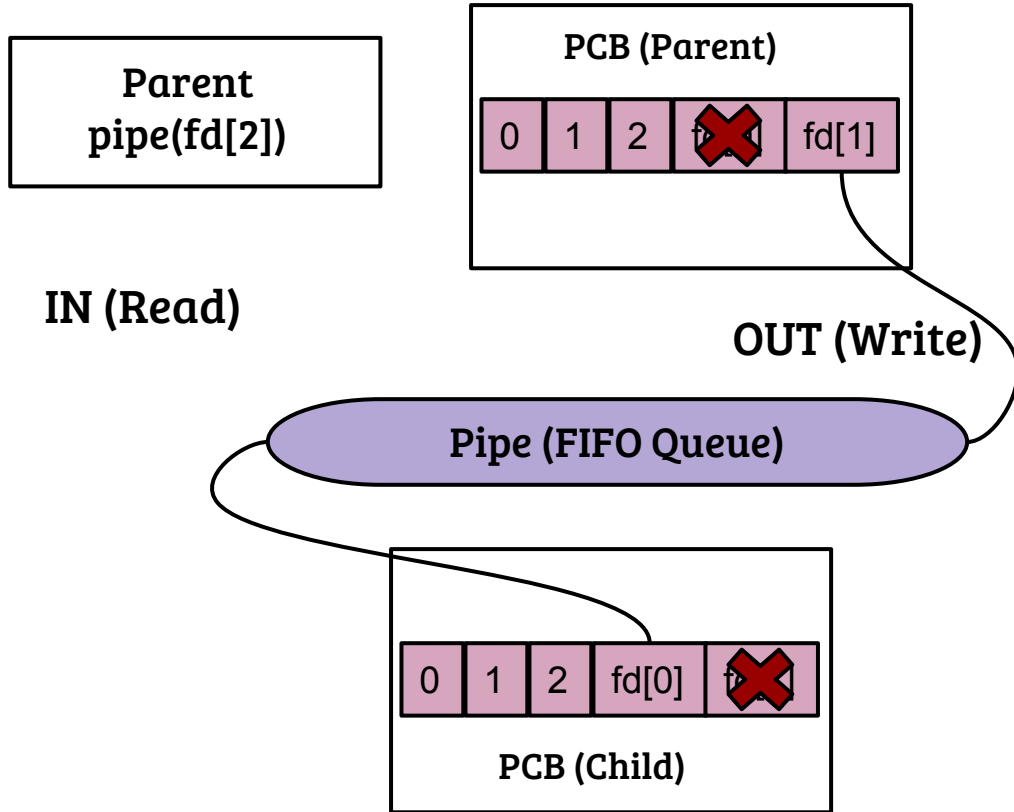
- `pipe()` takes array of two FDs as input
- `fd[0]` is the read end of the pipe
- `fd[1]` is the write end of the pipe
- Implemented as a FIFO queue in OS

UNIX pipe() with fork()



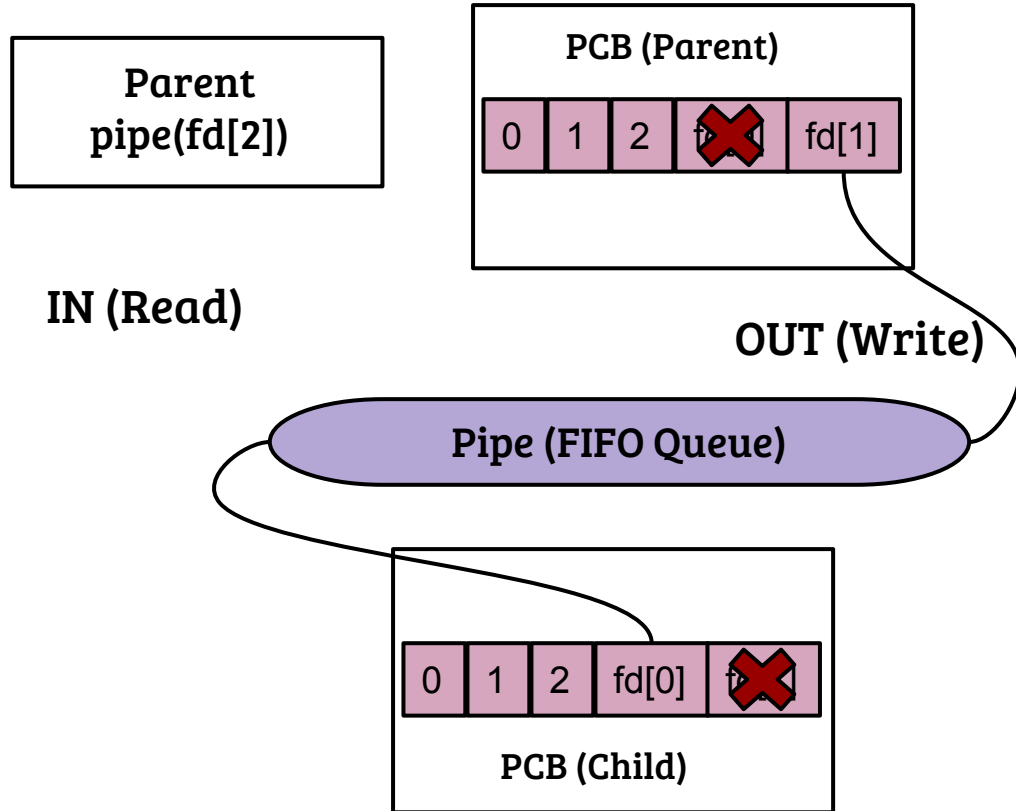
- `fork()` duplicates the file descriptors
- `close()` one end of the pipe, both in child and parent
- Result: a queue between parent and child

UNIX pipe() with fork()




- fork() duplicates the file descriptors
- close() one end of the pipe, both in child and parent
- Result: a queue between parent and child

Shell piping : pipe() + dup() + fork() + exec()



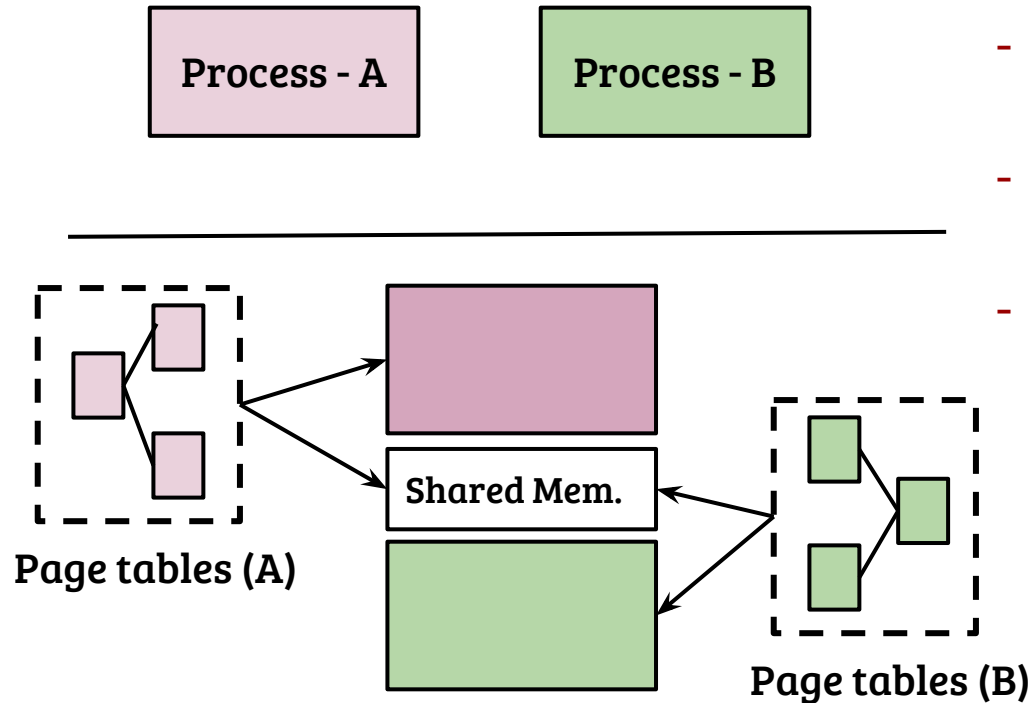
- pipe() followed by fork()
- exec("ls") after closing STDOUT and duping OUT fd of pipe
- exec("wc") after closing STDIN and duping IN fd of pipe
- Result: input of "wc" is connected to output of "ls"

Shared memory



- Shared memory made accessible using virtual memory
- How?

Shared memory

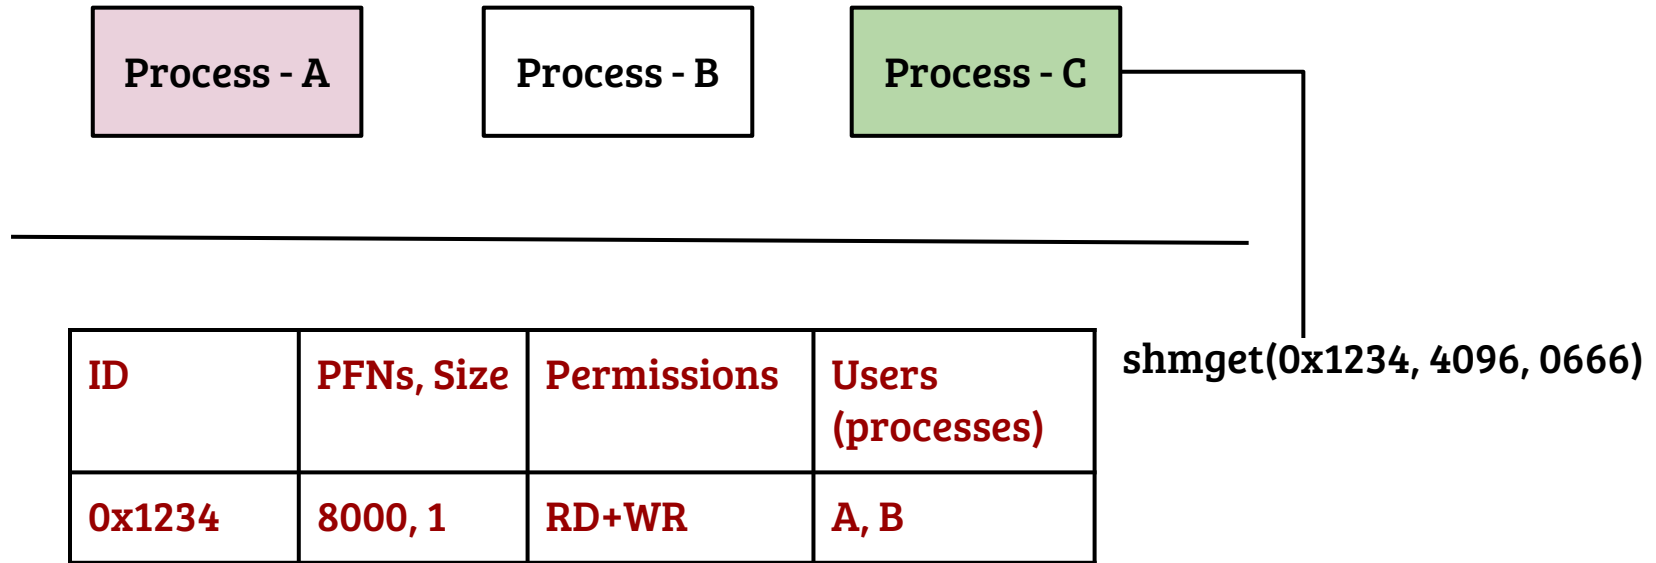


- A and B both map the shared region
- Is it required to be mapped to same VA in both processes?
- How shared memory regions are managed?

Shared memory design (API)

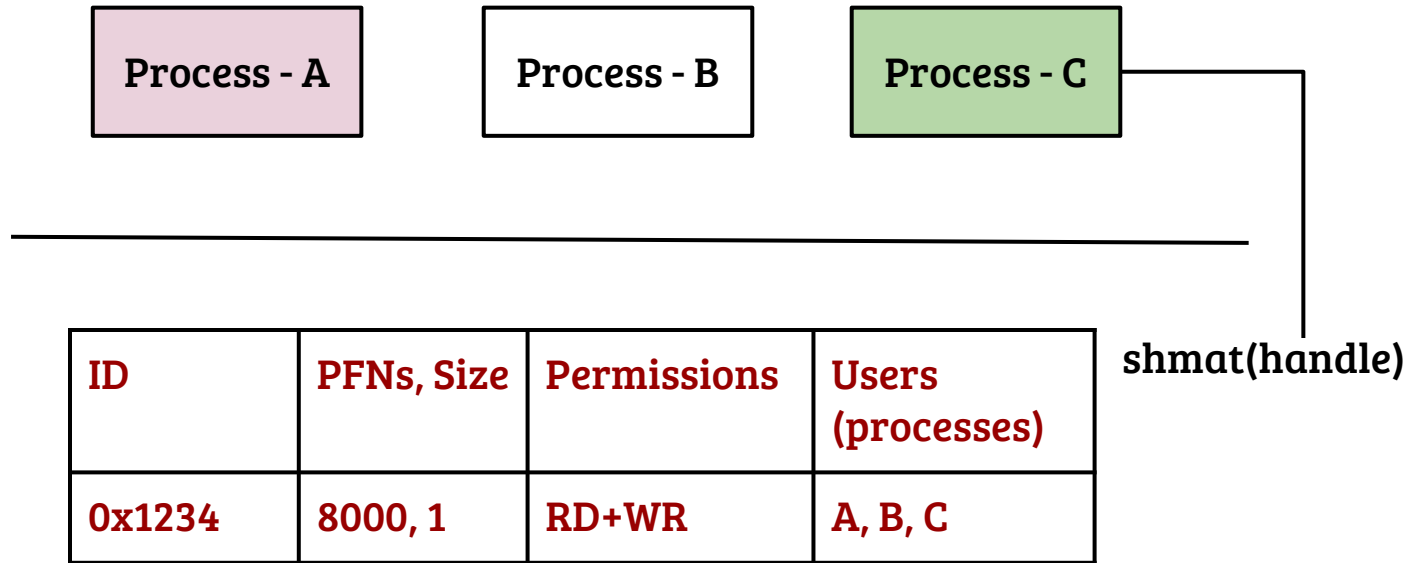
- How a shared memory region is created?
 - Must have a global identity
 - The OS maintains a list of shared regions
 - `handle = shmget (key, size, flags = IPC_CREAT|0666)`
- How any process gets an handle to a shared memory region?
 - Process must identify the shared region
 - The OS looks up in the list of shared regions
 - `handle = shmget (key, size, flags = 0666)`
- How any process gets the VA to the shared region?
 - Page table mapping inserted
 - `shmat(handle, address_hint, flags)`

Shared memory design (OS)



- OS looks up shared region in the global shared region table

Shared memory design (OS)



- OS creates V to P mapping for process C

Shared memory across fork and exec

- The child inherits the attached shared memory segments
- On `exec()`, all attached shared memory segments are detached
- On `exit()`, all attached shared memory segments are detached