

# Operating Systems

Semaphores, Classical problems

Debadatta Mishra, CSE, IITK

# Why mutual exclusion is not sufficient?

- Locking techniques allows exactly one thread to access the critical section
- Consider a scenario when at most  $K$  concurrent accesses are alright
- Example:
  - A finite array of size  $N$  accessed from a set of producer and consumer threads
  - At most  $N$  concurrent producers are allowed if array is empty
  - At most  $N$  concurrent consumers are allowed if array is full
- Using mutual exclusion techniques
  - Any access to array must take a lock, one thread allowed at a time, not very efficient
  - Does read-write locks help?

# Semaphores

```
typedef struct semaphore{
    int value;
    spinlock *LOCK;
    Queue *waitQ;
    ....
}sem_t; // Example

int wait (sem_t *s)
{
    s->value--;
    Wait if s->value <= 0
}

int post (sem_t *s)
{
    s->value++;
    wakeup one if one or more are waiting
}
```

- Generally, semaphores are initialized to a positive integer K
- Two operations: wait and post (other notations {wait, signal}, {P,V}, {down, up})

# Unix semaphores

```
#include <semaphore.h>

main( ){
    sem_t s;
    int K = 5;
    sem_init(&s, 0, K);
    sem_wait(&s);
    sem_post(&s);
}
```

- Can be used to synchronize threads of a single process if second argument is 0
- K is the value of the semaphore
- If K == 1, semaphore is called binary semaphores (mutex) as it can implement mutual exclusion
- lock → sem\_wait(&s)
- unlock → sem\_post(&s)

# Semaphore implementation (buggy #1)

```
wait(sem_t *s)
{
    while(s->value <= 0);
    s->value--;
}
post(sem_t *s)
{
    s->value++;
}
```

- What is wrong with this implementation?
- Atomic decrement and increment, will it help?
- Disable preemption?

# Semaphore implementation (inefficient)

```
wait(sem_t *s)
{
    lock(s->LOCK);
    while(s->value <= 0);
    s->value--;
    unlock(s->LOCK);
}
post(sem_t *s)
{
    atomic_inc(s->value);
}
```

- Busy wait, wastage of CPU cycles
- Why not put the context to sleep?

# Semaphore implementation (buggy#2)

```
wait(sem_t *s)                post(sem_t *s)
{                               {
    s->value--;                s->value++;
    if (s->value < 0){          if(s->value <= 0){
        insert_tail(s->waitQ, self);
        self->state = WAITING;
        schedule( );
    }
}
```

- Why is this not correct?
- Mutual exclusion is needed, but where exactly?

# Semaphore implementation (buggy#3)

```
wait(sem_t *s)
{
    lock(s->LOCK);
    s->value--;
    if (s->value < 0){
        insert_tail(s->waitQ, self);
        self->state = WAITING;
        schedule( );
    }
    unlock(s->LOCK);
}
```

```
post(sem_t *s)
{
    lock(s->LOCK);
    s->value++;
    if (s->value <= 0){
        p = remove_head(s->waitQ);
        p->state = READY;
    }
    unlock(s->LOCK);
}
```

- Why is this not correct?



# Semaphore implementation (buggy#4)

```
wait(sem_t *s)
{
    lock(s->LOCK);
    s->value--;
    if (s->value < 0){
        insert_tail(s->waitQ, self);
        self->state = WAITING;
        unlock(s->LOCK);
        schedule( );
    }
    unlock(s->LOCK);
}
```

```
post(sem_t *s)
{
    lock(s->LOCK);
    s->value++;
    if (s->value <= 0){
        p = remove_head(s->waitQ);
        p->state = READY;
    }
    unlock(s->LOCK);
}
```

- Failure Scenario?

# Semaphore implementation

```
wait(sem_t *s)
{
    lock(s->LOCK);
    s->value--;
    if (s->value < 0){
        insert_tail(s->waitQ, self);
        self->state = WAITING;
        unlock(s->LOCK);
        schedule( );
        return;
    }
    unlock(s->LOCK);
}
```

```
post(sem_t *s)
{
    lock(s->LOCK);
    s->value++;
    if (s->value <= 0){
        p = remove_head(s->waitQ);
        p->state = READY;
    }
    unlock(s->LOCK);
}
```

- Any assumptions for correctness?

# Semaphore usage example: wait for child

```
sem_init(&s, 0, 0);

child( ){
    .....
    post(s);
}

parent( ){
    create_child( );
    wait(s);
}
```

- Semaphore value initialized to zero
- If parent gets scheduled after the child creation, waits till child finished
- If child gets scheduled before parent, parent does not wait

# Semaphore usage example: ordering

```
A=0; B=0;
```

```
Th0 {  
    A = 1;  
    printf("%d\n", B);  
}
```

```
Th1 {  
    B=1;  
    printf("%d\n", A);  
}
```

- What are the possible outputs?
- How to ensure output (A = 1, B = 1)?

# Semaphore usage example: ordering

```
// s1->value = 0
A=0; B=0;

Th0{
    A = 1;
    wait(s1);
    printf("%d\n", B);
}
Th1{
    B=1;
    post(s1);
    printf("%d\n", A);
}
```

- What are the possible outputs?
- How to ensure output (A = 1, B = 1)?

# Ordering with two semaphores

```
// s1->value = 0, s2->value = 0  
A=0; B=0;
```

```
Th0{  
    A = 1;  
    post(s1);  
    wait(s2);  
    printf("%d\n", B);  
}
```

```
// s1->value = 0, s2->value = 0  
A=0; B=0;
```

```
Th1{  
    B=1;  
    wait(s1);  
    post(s2);  
    printf("%d\n", A);  
}
```

# Producer-consumer problem

- A buffer of size  $N$ , one or more producers and consumers
- Producer adds an element to the buffer
- Consumer extracts an element from the buffer
- Example: A multithreaded web server, network device queue etc.
- Solution using semaphores?

# Producer-consumer problem (buggy #1)

```
item A[N], sem_t empty = {N}, used = {0}, pctr = 0, cctr = 0;
```

```
produce(item x)
{
    wait(empty);
    A[pctr] = x;
    pctr = (pctr + 1) % N;
    post(used);
}
```

```
consume()
{
    item x;
    wait(used);
    x = A[cctr];
    cctr = (cctr + 1) % N;
    post(empty);
    return x;
}
```

- Both produce and consume can be called concurrently
- What is the problem?



# Producer-consumer problem (buggy #2)

```
item A[N], sem_t empty = {N}, used = {0}, pctr = 0, cctr = 0, mutex M;
```

```
produce(item x)
{
    lock(M);
    wait(empty);
    A[pctr] = x;
    pctr = (pctr + 1) % N;
    post(used);
    unlock(M);
}
```

```
consume()
{
    item x;
    lock(M)
    wait(used);
    x = A[cctr];
    cctr = (cctr + 1) % N;
    post(empty);
    unlock(M);
    return x;
}
```

- What is the problem?

# Producer-consumer problem

```
item A[N], sem_t empty = {N}, used = {0}, pctr = 0, cctr = 0, mutex M;
```

```
produce(item x)
{
    wait(empty);
    lock(M);
    A[pctr] = x;
    pctr = (pctr + 1) % N;
    unlock(M);
    post(used);
}
```

```
consume()
{
    item x;
    wait(used);
    lock(M);
    x = A[cctr];
    cctr = (cctr + 1) % N;
    unlock(M);
    post(empty);
    return x;
}
```

- What if separate mutex is used for producer and consumer?

# Producer-consumer problem (buggy?)

```
item A[N], sem_t empty = {N}, used = {0}, pctr = 0, cctr = 0, mutex P, C;
```

```
produce(item x)
{
    wait(empty);
    lock(P);
    A[pctr] = x;
    pctr = (pctr + 1) % N;
    unlock(P);
    post(used);
}
```

```
consume()
{
    item x;
    wait(used);
    lock(C);
    x = A[cctr];
    cctr = (cctr + 1) % N;
    unlock(C);
    post(empty);
    return x;
}
```

- Is there any issue?

# Producer-consumer variant

```
produce(item x)
{
    wait(empty);
    lock(P);
    copy(A[pctr], x);
    pctr = (pctr + 1) % N;
    unlock(P);
    post(used);
}
```

```
consume()
{
    item x;
    wait(used);
    lock(C);
    copy(x, A[cctr]);
    cctr = (cctr + 1) % N;
    unlock(C);
    post(empty);
    return x;
}
```

- Copy inside the mutex not very efficient, can it be moved outside?

# Producer-consumer variant (buggy?)

```
produce(item x)
{
    int lpos;
    wait(empty);
    lock(P);
    lpos = pctr;
    pctr = (pctr + 1) % N;
    unlock(P);
    copy(A[lpos], x);
    post(used);
}
```

```
consume()
{
    item x; int lpos;
    wait(used);
    lock(C);
    lpos = cctr;
    cctr = (cctr + 1) % N;
    unlock(C);
    copy(x, A[lpos]);
    post(empty);
}
```

- Does this code work?

# Producer-consumer variant

```
produce(item x)
{
    .....
    pctr = (pctr + 1) % N;
    lock(A[lpos].lock);
    unlock(P);
    copy(A[lpos], x);
    unlock(A[lpos].lock);
    post(used);
}
```

```
consume()
{
    .....
    cctr = (cctr + 1) % N;
    lock(A[lpos].lock);
    unlock(C);
    post(empty);
    copy(x, A[lpos]);
    unlock(A[lpos].lock);
}
```

- Why does this code work?
- Any benefits?

# Condition variables

- *pthread\_cond\_wait (cond, mutex)*: Atomically releases the mutex and waits on a condition variable. Imp: *Does not perform any condition check*. Resumes execution holding the lock when *pthread\_cond\_signal( )* is invoked
- *pthread\_cond\_signal (cond)*: Wakes up a waiting thread on condition *cond*, Ideally called holding the mutex.

# Example

```
cond_t C; mutex M;  BOOL condition;
```

```
Th1()  
{  
    while(1){  
        .....  
        lock(M);  
        while(condition != true)  
            cond_wait (C, M);  
        unlock(M);  
        .....  
    }  
}
```

```
Th2()  
{  
    while(1){  
        .....  
        lock(M);  
        condition = true;  
        cond_signal(C);  
        unlock(M);  
        .....  
    }  
}
```

- What if Th2( ) does not acquire lock?
- What if cond\_signal( ) called after unlock(M)?



# An example implementation (src: pthread man page)

```
cond_t { mutex CM, int val, waitQ};
```

```
cond_wait(cond_t C, mutex M)
{
    int value = C->value;
    unlock(M);
    lock(C->CM);
    If (value == C->value){
        addQ(C->waitQ, self);
        unlock(C->CM);
        schedule( );
    }else {unlock(C->CM);}
    lock(M);
}
```

```
cond_signal(cond_t C)
{
    lock(C->CM);
    C->value++;
    if(!empty(C->waitQ)){
        p=pickQ(C->waitQ);
        p->state = running;
    }
    unlock(C->CM);
}
```

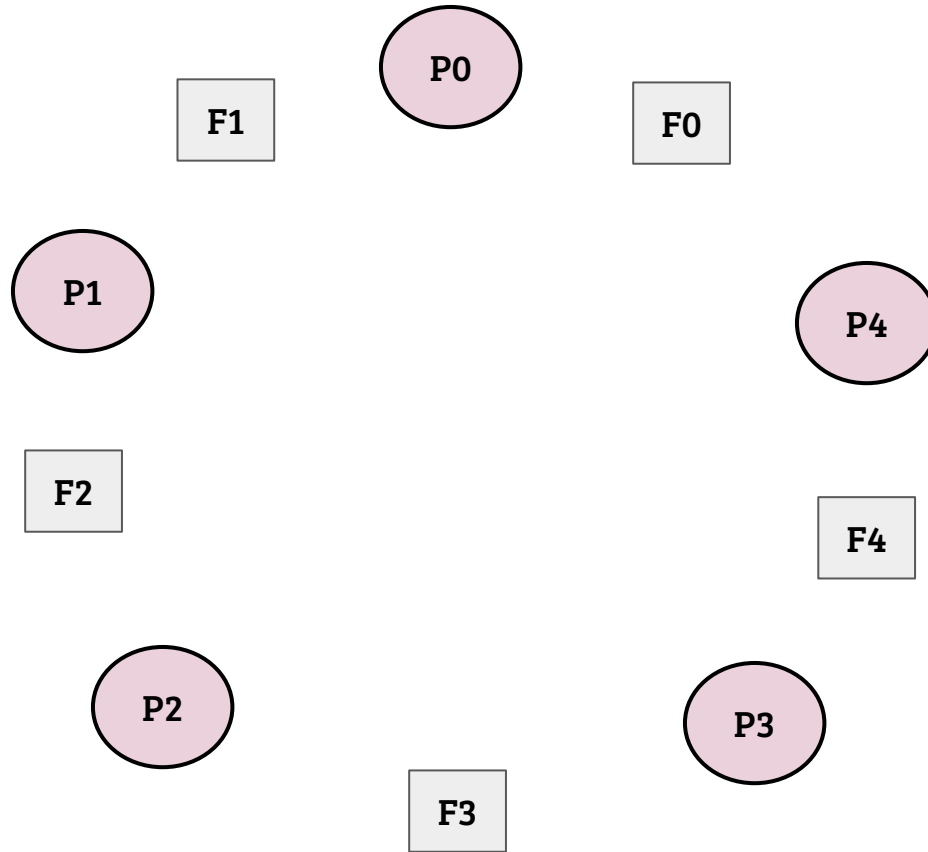
# Deadlock: example

```
ACCOUNT {  
    long ID;  
    mutex M;  
    long balance;  
    ...  
};
```

```
Txn_transfer(ACCOUNT *S, ACCOUNT *D, $)  
{  
    lock (S->M);  
    lock (D->M);  
    .....  
    .....  
    unlock (D->M);  
    unlock(S->M);  
}
```

- Where is the deadlock?
- Example: {T1: Txn\_transfer(&CSE, & IITK);}, {T2: Txn\_transfer(&IITK, &CSE);}
- Solution?

# Dining philosophers



```
P(ID )  
{  
    while (1) {  
        think( );  
        getforks( );  
        eat( );  
        putforks( );  
    }  
}
```

# Conditions for deadlock

- Mutual exclusion → Exclusive access to the resource
- Hold-and-wait → Hold one lock and wait for other
- No resource preemption → Locks can not be forcibly removed from threads holding them
- Circular wait → A cycle of threads requesting locks held by others.  
Specifically, a cycle in the directed graph  $G(V, E)$  where  $V$  is the set of processes,  $(v1, v2) \in E$  if  $v1$  is waiting for a lock held by  $v2$

All of the above conditions should be satisfied for a deadlock to occur