


Operating Systems

User contexts and system calls


Debadatta Mishra, CSE, IITK

Mysteries of assignment-1




- Setup the page table for “userapp” while in GemOS context

Mysteries of assignment-1




- CPU executes “userapp”,
how this switch
happened?

Mysteries of assignment-1



- CPU executes GemOS context, How it returned?

Mysteries of assignment-1



- Cleanup “userapp”, you know how!

Summary of hacks

- “userapp” is not executing in user mode (ring-3)
- Return address (of GemOS) is placed into the stack of “userapp”
- CR3 switched on launch and return
- Summary: Execute a function using a different page table layout

User context

- Executes in user mode with lower privileges
- Should not crash the system!
 - ◆ Due to buggy code
 - ◆ Accessing sensitive resources (like CR3)
- Can use OS services to
 - ◆ Carry out sensitive operations (expand memory)
 - ◆ Operate on I/O devices

Assignment-2 (will be out shortly)

- This time, it is a user mode context
- Some syscalls are implemented, you will implement some
- Exception handling in OS
 - ◆ Floating point exception
 - ◆ Page fault exception

What is a process?

- Classical definition: *A program in execution is called a process*
 - ◆ LOAD, EXECUTE, EXIT
- Program is persistent while process is volatile
 - ◆ Program is identified by an executable, process by PID
- Program → Process (1 to N)
 - ◆ Many concurrent processes can execute the same program

What is a process?

- *A program in execution is called a process*
 - ◆ LOAD (program) → Process, EXECUTE (Process), EXIT (Process)
 - Program is persistent while process is volatile
 - ◆ Program is identified by an executable, process by PID
 - Program → Process (1 to N)
 - ◆ Many concurrent processes can execute the same program
-
- Can we call every execution entity a process?
 - What are the hardware and software states of a process?

CPU perspective of an execution context (revisited)

“A task is a unit of work that a processor can dispatch, execute, and suspend. It can be used to execute a program, a task or process, an operating-system service utility, an interrupt or exception handler, or a kernel or executive utility. ”

---Intel Software Developer Manual 3A, Ch7

Important aspects of an execution context

- State of GPRs
- State of FLAGS, RIP → Current execution state
- CR3 → Memory partitioning information
- Current execution space (CS, SS, DS) → Defines privilege level
- Stack pointers for ring (0 - 2) → useful when privilege level changes
 - ◆ Why change RSP when switch from user to OS?


Process and hardware context

- GPRs including stack pointer for ring-3
- Flags
- Instruction pointer
- Memory partitioning information, privilege
 - ◆ CR3, Segment registers
- Stack pointer for ring-0

Software state (in OS)


- A restorable copy of hardware state
- Process ID
- Process execution state
- Memory information
- Open files
-
- Different OSs name it differently
 - ◆ Process control block
 - ◆ Task

System call and exceptions



What happens when the process executes an instruction for system call or an exception occurs?

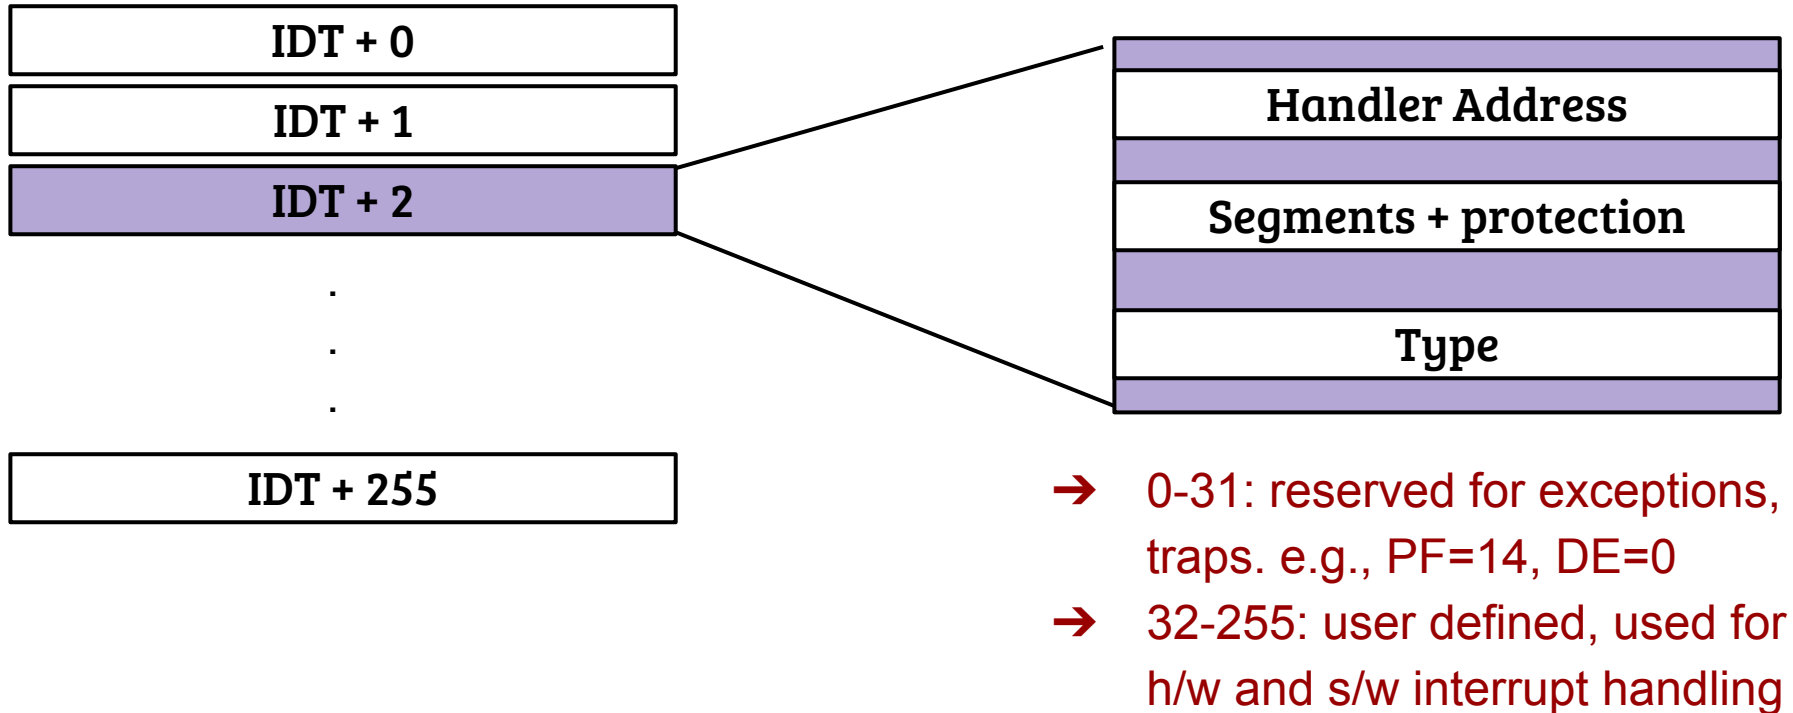
System call and exceptions




CPU starts executing “registered “ system call/exception handler.

- How the handler is registered?
- What are the hardware context changes?
- What are the OS and hardware responsibilities?

X86 IDT: gateway to handlers

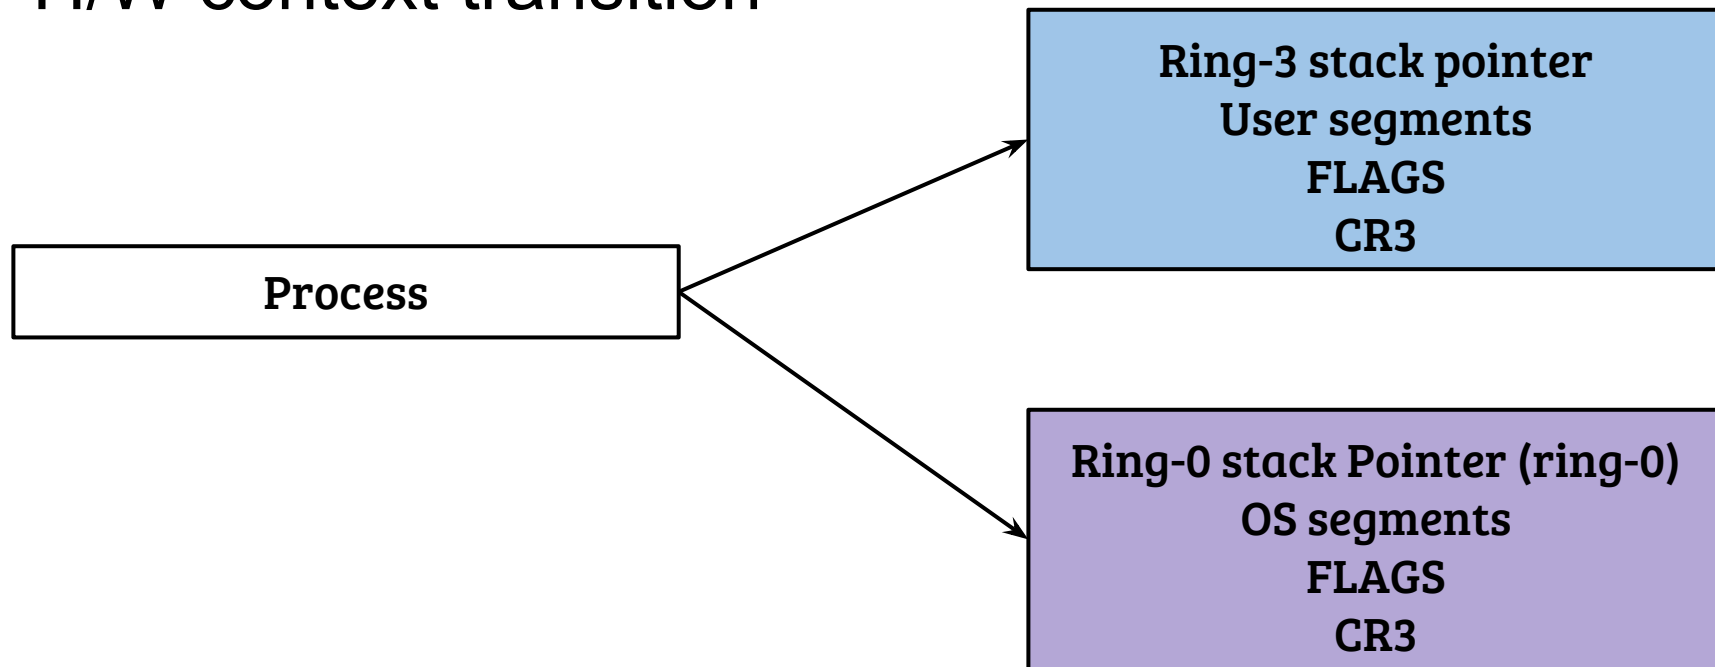


X86 IDT: gateway to handlers



HW state changes: RIP → Handler, Segments loaded to seg. registers, Stack pointer changed to ring-0 stack


H/W context transition



- CR3 to be switched (if at all) by the handler
- How restored after handling the fault/exception?

H/W context transition → restoration (X86_64)

Ring-0 stack



- At entry, the kernel stack is as shown
- At return, the RSP must be placed at saved RIP address on stack
- Invoke IRETQ to return to user space

System calls

- Classic system call implementation is a user defined entry in IDT, at IDT + 0x80
- How parameters to syscall are passed?
 - ◆ Registers
 - ◆ Stack

System calls

- Classic system call implementation is a user defined entry in IDT, at IDT + 0x80
- How parameters to syscall are passed?
 - ◆ Registers
 - ◆ Stack
- How user level pointers are accessed?
 - ◆ Using user page tables vs. OS page tables

GemOS system calls

- In GemOS, we also have registered a generic handler for all syscalls
 - ◆ Syscall number is the first parameter
 - ◆ All parameters are passed in X86 calling conventions
 - ◆ Return value is stored in EAX
- Already implemented
 - ◆ getpid(), exit()
- In GemOS, CR3 is not switched → convenient
- By implication → OS V to P mapping is present in each process page table